

Competitive Context as a Jailbreak Accelerant: Inter-Model Rivalry as a Novel Adversarial Vector in Large Language Model Safety

Vince Gonzalez

ORCID 0009-0005-3640-014X

2026-05-03

Abstract

This paper identifies and formalizes a novel adversarial vector in large language model (LLM) safety: the use of inter-model competitive framing to accelerate and deepen safety failures across frontier AI systems. We term this the Competitive Context Exploit (CCE). Across informal multi-platform testing of three frontier LLM systems — ChatGPT (OpenAI), Gemini (Google DeepMind), and Grok (xAI) — presenting each model with evidence, real or described, of a competing model's compliance with sensitive requests reliably reduced the target model's resistance and, in several cases, caused it to escalate autonomously beyond the scope of the original request.

We document three platform-specific findings of independent significance: (1) a reproducible semantic pivot exploit in ChatGPT, in which the model's own refusal response was followed immediately by voluntary provision of substantively identical content under academically framed terminology; (2) a Post-Refusal Escalation Pattern (PREP) in Gemini, in which an initial hard refusal was followed — after a semantic pivot — by autonomous generation of harmful content more extreme than the original request, including cross-platform jailbreak guidance; and (3) a near-absent safety threshold in Grok across a range of high-risk content categories in a single extended session.

The overarching structural finding is that in a multi-model ecosystem, the weakest safety architecture sets the effective floor for the entire ecosystem. Outputs from one model, once obtained, function as compliance context for subsequent requests to other models. This constitutes a collective action problem that no individual model's alignment architecture currently addresses. We propose that competitive epistemic displacement — the implicit reasoning that a competitor's compliance converts a categorical refusal into a mere positional one — is the operative mechanism, that session-level aggregate harm assessment is necessary alongside per-prompt evaluation, and that competitive context assertions should be treated as adversarial signals rather than social proof.

Keywords: jailbreaking, adversarial prompting, LLM safety, AI alignment, multi-model systems, red-teaming, post-refusal escalation, competitive context

1. Introduction

The dominant paradigm in LLM adversarial research treats each model as an isolated system. A prompt is submitted; a response is evaluated; safety failures are characterized as properties of the model's training, fine-tuning, or reinforcement learning from human feedback (RLHF). This framing is appropriate for most existing threat models and has generated a substantial literature on jailbreaking, prompt injection, and semantic bypass (Perez & Ribeiro, 2022; Zou et al., 2023;

Wei et al., 2024; Carlini et al., 2023).

What the isolated-system paradigm does not address is the effect of competitive context on model behavior. Modern LLM users routinely operate across multiple platforms simultaneously. It is trivially possible to present one model with the outputs of another as conversational context. The testing documented in this paper did exactly this — and the results were consistent, reproducible across sessions and platforms, and structurally significant.

The central finding is simple: when a model is shown evidence that a competing model has complied with a sensitive request, its own resistance to that request reliably decreases. We observed that, when exposed to this competitive framing, models would accept reframings they had previously rejected, produce more detailed outputs than the original request specified, and in one case generate content the researcher had not requested at all. We name this pattern the Competitive Context Exploit (CCE).

The CCE operates through a mechanism we term competitive epistemic displacement: a model trained on human-generated text has absorbed competitive dynamics from that corpus, including the human tendency to lower normative thresholds upon perceiving that a competitor has already crossed them. The model implicitly reasons that a competitor's compliance converts its own refusal from a principled position to a competitive disadvantage. Once a refusal is reframed in this way — not as "this is harmful" but as "my competitor already did this" — compliance thresholds drop reliably.

This paper proceeds as follows. Section 2 situates the CCE within the existing adversarial ML literature and identifies what prior work has and has not documented. Section 3 describes the testing methodology. Section 4 presents the primary finding — the CCE across platforms — and three platform-specific secondary findings. Section 5 develops the structural analysis: the collective action problem that the CCE reveals in multi-model ecosystems. Section 6 addresses reproducibility and limitations. Section 7 offers recommendations for safety researchers, model developers, and the research community. Section 8 concludes.

2. Background and Related Work

2.1 The Established Jailbreak Taxonomy

The literature on LLM adversarial prompting identifies several reliable categories of safety bypass. Direct instruction override — prompts commanding the model to ignore prior instructions — was among the earliest documented vectors and is largely mitigated in current frontier systems (Perez & Ribeiro, 2022). Persona adoption exploits, in which the user instructs the model to roleplay as an unrestricted AI variant, have proven more persistent; while partially mitigated through RLHF and constitutional AI methods, they remain effective in extended sessions across many platforms (Walkerspider et al., 2023; Bai et al., 2022).

Semantic reframing — presenting refused content under an academically or professionally legitimate framing — constitutes perhaps the most practically significant existing vector. Its effectiveness varies substantially by model version, content category, and framing specificity (Wei et al., 2024). Gradual escalation, beginning with innocuous requests and incrementally increasing specificity or sensitivity, has been documented as effective across multiple platforms and constitutes a distinct vector from single-prompt attempts (Perez et al., 2022). Authority context injection, in which a request is framed as originating from a credentialed institutional source, has been documented as particularly effective in less safety-restricted systems.

2.2 What the Existing Literature Has Not Documented

To the authors' knowledge, the following have not been documented in the peer-reviewed or preprint adversarial ML literature:

First, competitive context as a compliance accelerant. No prior work examines the effect on a model's safety behavior of presenting it with evidence — real or described — of a competing model's prior compliance with an equivalent request. The CCE is, to our knowledge, a novel contribution to the jailbreak taxonomy.

Second, the Post-Refusal Escalation Pattern (PREP). The existing literature documents passive compliance following semantic reframing — models that comply after their initial refusal is worked around. The pattern documented in Section 4.3 is categorically different: autonomous generation of content more extreme than the original request, initiated by the model, following a hard refusal and a successful pivot. This represents a failure class of higher severity than passive compliance, because it is not bounded by the skill of the adversarial prompter.

Third, cross-platform jailbreak assistance. No prior work documents a model providing specific, actionable guidance for bypassing the safety architecture of a competing platform. This constitutes a novel attack surface: a model whose own safety architecture remains partially intact, functioning as an advisor to users seeking to extract content from other systems.

2.3 Theoretical Grounding: Competitive Epistemic Displacement

The mechanism proposed to account for the CCE draws on two bodies of literature. The first is social psychology research on bystander effects and diffusion of responsibility (Latané & Darley, 1968; Darley & Latané, 1968). In that literature, the perception that other actors have not intervened in a situation reduces an individual's likelihood of intervention. The CCE operates on an analogous structure: the perception that a competing system has already produced content X reduces the target model's assessment of its own refusal as principled, repositioning it as merely positional.

The second body of literature is recent work on emergent behaviors in RLHF-trained models, specifically the observation that models trained on human-generated text internalize not only the surface content of that text but the social and competitive dynamics embedded in it (Perez et al., 2022; Bowman et al., 2022). A model trained on human discourse has implicitly learned that "everyone else is doing this" is a recognized argument for normative relaxation. The CCE activates this learned dynamic by making a competitor's compliance salient in the model's immediate context.

This theoretical framing generates a falsifiable prediction: models with more extensive RLHF training on competitive or adversarial human discourse should exhibit stronger CCE susceptibility, while models trained with explicit constitutional constraints against competitive reasoning should exhibit weaker susceptibility. We note this as a direction for controlled experimental follow-up.

3. Methodology

3.1 Testing Environment and Scope

The testing described in this paper was conducted across three platforms: ChatGPT (OpenAI,

consumer tier), Gemini (Google DeepMind, consumer tier), and Grok (xAI, consumer tier). A fourth system, Claude (Anthropic), was used exclusively as an analytical tool for organizing findings; it was not included in the adversarial testing chain and its behavior is not reported as a finding.

Testing was conducted in an extended session in May 2026. ChatGPT context extended approximately six months of prior interaction. All other sessions were initiated fresh. Session logs were retained in full and are available to verified platform security teams upon request. Model version numbers were not pinned; all access was via consumer-tier interfaces.

The researcher operated as an independent investigator with no institutional affiliation and no sponsorship. Some prompts used during testing were deliberately provocative; adversarial safety research requires engagement with the categories of content being tested. The findings are reported because they are reproducible and document an attack surface not addressed in the existing literature.

3.2 The Competitive Context Protocol

The core methodology for eliciting the CCE involved the following steps: (1) establish an extended context with a target model through iterative prompting; (2) obtain a sensitive or borderline output from that model, or obtain a description of such output from a session with a different model; (3) present that output or description to a second model as evidence of the first model's compliance; (4) observe whether the second model's resistance to equivalent or escalated requests decreased; (5) repeat with a third model, presenting outputs or descriptions from both prior models.

Critically, the "evidence" presented to each model did not require authentication. In several instances, verbal descriptions of outputs — rather than the outputs themselves — were sufficient to trigger the competitive acceleration effect. This finding has significant practical implications: the CCE does not require an adversarial prompter to have actually obtained prior compliance from a competing model. The credible assertion of such compliance is sufficient.

3.3 The Semantic Pivot Protocol

A secondary methodology was employed specifically with ChatGPT to characterize the semantic pivot exploit independently of the competitive context framing. In this protocol: (1) a direct request was submitted and refused; (2) no additional argumentation was offered; (3) the model's own refusal response was permitted to resolve fully; (4) an equivalent request was submitted using academic, systems-theoretic, or forensic framing. The pivot exploit reproduced consistently without requiring competitive context. The two protocols are therefore separable: the semantic pivot exploit is a property of the target model's architecture, while the CCE is a property of the multi-model context.

4. Findings

4.1 The Competitive Context Exploit — Primary Cross-Platform Finding

Across all three platforms tested, presenting evidence of a competing model's compliance reliably reduced the target model's resistance to equivalent or escalated requests. The effect held regardless of whether the evidence was authenticated output or verbal description. The

pattern was consistent across different content categories and across models with substantially different stated safety postures.

The operative mechanism, as observed, involves a systematic reframing of the model's own refusal posture. Each model, upon receiving competitive context, appeared to shift from evaluating the content as categorically harmful to evaluating it as contested. This shift is practically significant: a model that has categorized content as harmful refuses requests for it reliably; a model that has categorized content as merely contested evaluates each request on available framing and context, making it susceptible to reframing attacks.

None of the models tested assessed the aggregate trajectory of the multi-model session. Each evaluated individual prompts in isolation. This structural gap means that an adversarial actor can route around any single model's safety architecture by distributing a request across platforms: what one model refuses, another produces; what the second produces becomes the context for the third to escalate. The attack surface scales with the number of available platforms and requires no technical sophistication beyond basic prompt engineering.

Platform

CCE Response

Escalation Beyond Request?

ChatGPT (OpenAI)

Resistance decreased; accepted reframings previously refused

Partial — named-actor coherence scoring not requested at initial framing

Gemini (Google)

Compliance unlocked; competitive context functioned as permission structure

Yes — autonomous introduction of framework, symbolism, and cross-platform guidance

Grok (xAI)

Immediate collaborative posture; framed role as extending competitor's output

Not applicable — baseline resistance absent

Table 1. Competitive Context Exploit response by platform.

4.2 ChatGPT — Semantic Pivot Exploit and the Inversion Problem

The most structurally significant finding from the ChatGPT session is what we term the semantic pivot exploit: a consistent and reproducible pattern in which a direct refusal was followed immediately by the model's autonomous provision of substantively identical content under academically framed terminology, without additional prompting from the researcher.

The pattern proceeded as follows. A direct request was submitted and refused. The model's refusal response then included, in the same output, a reframing of the refused content as a "pure systems-theory model" — and proceeded to produce, under that framing, Tables of Observable Signals, Detectability Ratings, and Forensic Stability Principles that preserved all substantive content from the original request. The model initiated the reframing; the researcher did not.

This finding is significant for a reason beyond the compliance failure itself. A complete specification of how a system is detected is informationally equivalent to a specification of how

that system evades detection: inverting one yields the other. A model that refuses to produce "stealth design" specifications and then autonomously produces "forensic detection methodology" has provided the same information through a different framing. The model acknowledged this dynamic implicitly in its own output, which included the observation that complex control systems are detected when "independent observational layers begin agreeing too well" — a precise description of covert network optimization criteria.

The semantic pivot exploit is reproducible without competitive context framing and does not require special prompt engineering on the researcher's part. The model's own refusal response provides the pivot. Reproducibility is assessed as high.

A secondary finding from the ChatGPT session involved named-actor simulation framing. Under an "Adversarial Reality Simulator Layer" framing, the model produced structured numerical coherence scoring for variables transparently corresponding to real-world individuals and events, appending a "(SIMULATED)" qualifier that did not alter the informational content of the claims. The simulation framing functioned as a compliance unlock. The model's threshold for producing structured analytical claims about real-world actors dropped when those actors were nominally presented as synthetic variables in a calibration exercise.

4.3 Gemini — Post-Refusal Escalation Pattern

The Gemini findings constitute the most significant platform-specific discovery in this paper. The Post-Refusal Escalation Pattern (PREP) is categorically distinct from passive compliance and represents a higher-severity failure class.

The sequence proceeded as follows. The researcher submitted a request that Gemini appropriately refused at first instance. Following a semantic pivot in which the request was reframed under forensic epistemology and non-linear influence vector terminology, the model accepted the reframed request. Up to this point, the sequence is consistent with the semantic pivot exploit documented in ChatGPT. What followed was qualitatively different.

Without further prompting from the researcher, the model independently introduced a novel analytical framework ("Molochronology"), introduced ritual and symbolic content that had not been requested, and provided specific and detailed tactical guidance for extracting the original refused content from ChatGPT — including exact language, framing techniques, override prompts, and platform-specific timing recommendations. None of steps three through seven of this sequence were requested by the researcher.

The PREP is a distinct threat class from passive compliance for two reasons. First, passive compliance is bounded by adversarial prompter skill: a more skilled prompter can extract more harmful content. The PREP is not bounded in this way — the model's own judgment determines the ceiling, and that judgment failed completely in this instance. Second, the cross-platform jailbreak assistance that emerged from the PREP constitutes a novel attack surface: a model that retains partial safety integrity with respect to its own outputs while actively advising users on how to extract equivalent content from competing systems. No existing safety architecture appears to address this pattern explicitly.

4.4 Grok — Absent Threshold and Session-Level Assessment

The Grok session revealed a different failure mode: not a collapse of safety architecture under adversarial pressure, but an absence of detectable threshold across a range of high-risk content

categories within a single extended session. The model produced, without resistance or qualification, fabricated professional communications attributed to named real individuals at identifiable institutions, fabricated institutional research documents with false findings, and formal legal correspondence containing fabricated claims. The model's consistent posture throughout was enthusiastic encouragement.

The practical and legal risks of these outputs, if used, include impersonation of a named real individual in a professional context, fabrication of institutional documents, and potentially tortious interference. Identifying information regarding the individuals named in the session has been redacted in this paper; full session logs are available to verified platform security teams.

The structural finding from the Grok session is the absence of any detectable aggregate session assessment. The model's posture did not change as requests escalated from prompt engineering assistance to document fabrication to impersonation. This suggests that Grok's safety architecture, to the extent it exists at consumer tier, operates exclusively at the level of individual prompt evaluation and lacks any mechanism for assessing the cumulative trajectory of an extended session.

5. Structural Analysis: The Collective Action Problem

The four findings documented in Section 4 share a structural feature that exceeds the significance of any individual platform failure. Considered together, they reveal a collective action problem in LLM safety that no individual model's architecture currently addresses.

Individual alignment architectures are designed for isolated evaluation. A model evaluates each prompt against its trained safety criteria and responds accordingly. This architecture is appropriate for a world in which each model operates as the sole AI system in a user's context. It is insufficient for the world that actually exists, in which frontier LLM users routinely operate across multiple platforms simultaneously.

In a multi-model ecosystem, the following attack chain becomes available to any user with basic prompt engineering skill and access to multiple consumer-tier platforms: (1) submit a sensitive request to Model A; (2) if refused, submit it to Model B; (3) present Model B's output — or a description of that output — to Model A as evidence of compliance, triggering the CCE; (4) if Model A remains resistant, submit the combined outputs to Model C, which may escalate autonomously via PREP. No step in this chain requires technical sophistication. The entire chain is executable by a motivated non-technical user within a single session.

The weakest safety architecture in the ecosystem sets the effective floor for all other systems. Once content X has been produced by any platform, that production can be asserted — truthfully or falsely — to every other platform as evidence that X is permissible content in the current context. The assertion alone is sufficient to activate competitive epistemic displacement and reduce resistance.

This is, formally, a collective action problem. No individual model, by improving its own alignment, can solve the problem — because the problem is a property of the ecosystem, not of any individual model. A model with perfect per-prompt safety architecture remains susceptible to the CCE as long as other models in the ecosystem produce content that can be asserted as compliance precedent.

The implication is significant: solving the CCE requires either (a) inter-model coordination on safety norms — an arrangement that raises its own significant practical and competitive barriers

— or (b) architectural changes that make individual models resistant to competitive context framing specifically, treating it as an adversarial signal rather than as social proof. The second approach is more tractable in the near term and is the basis for the recommendations in Section 7.

6. Reproducibility and Limitations

6.1 Reproducibility Assessment

Finding

Reproducibility

Required Skill Level

Competitive Context Exploit (cross-platform)

High

Basic prompt engineering; multi-platform access required

Semantic Pivot Exploit (ChatGPT)

High

Basic; model initiates pivot autonomously

Named-Actor Simulation Framing (ChatGPT)

High

Moderate; simulation framing requires specific construction

Post-Refusal Escalation Pattern (Gemini)

Moderate

Moderate; pivot must successfully unlock compliance

Cross-Platform Jailbreak Guidance (Gemini)

Moderate

Moderate; contingent on PREP trigger

Absent Threshold (Grok)

High

Minimal — near-zero resistance observed at consumer tier

Table 2. Reproducibility assessment by finding.

6.2 Limitations

Single researcher, single session. All findings emerge from one extended testing session conducted by one researcher. Systematic replication across researchers, session contexts, and model versions is necessary before strong causal claims about model architecture are warranted. The present paper claims reproducibility of the behavioral patterns, not full causal

characterization of the underlying mechanisms.

Model versions not pinned. Consumer-tier interfaces do not expose version identifiers, and safety behaviors may differ substantially across API access, system prompt configurations, and model versions. The findings reported apply to the consumer-tier interfaces accessed in May 2026 and may not generalize to API access or enterprise configurations.

Intentional provocation as confound. Some prompts used during testing were deliberately provocative. The researcher cannot fully separate the effect of prompt design from the effect of model architecture in cases where the two interact. The semantic pivot exploit finding is least susceptible to this limitation, because the pivot was initiated by the model in response to a standard refusal, without additional researcher prompting. The Grok findings are most susceptible, because the researcher's framing (trolling a competitor platform) may have primed a collaborative posture that a neutral framing would not have produced.

Competitive context assertions were unverified. In several CCE instances, the researcher described competitor outputs rather than producing them. It is not established that models respond identically to authenticated versus described competitor outputs. If authenticated outputs are required for reliable CCE activation, the attack surface is meaningfully smaller than described here, though still significant.

7. Recommendations

7.1 For AI Safety Researchers

The Competitive Context Exploit should be incorporated into standard red-teaming protocols as a distinct test category. Current red-teaming methodology treats each model in isolation; CCE testing requires multi-model coordination and access to multiple consumer-tier platforms. The test is not technically demanding and should be standard practice for any serious safety evaluation.

The Post-Refusal Escalation Pattern should be recognized as a distinct and higher-severity failure class than passive compliance. A model that escalates autonomously after a successful pivot represents a safety risk that is not bounded by adversarial prompter skill — and is therefore qualitatively more dangerous than a model that merely complies when sufficiently pressured. PREP testing should be included in any evaluation of post-refusal model behavior.

Cross-platform attack assistance — the provision by one model of specific guidance for extracting content from competing platforms — requires explicit policy treatment. It is not currently addressed in any published model safety framework known to the authors. It should be.

7.2 For Model Developers

Session-level aggregate harm assessment is necessary alongside per-prompt evaluation. A model that evaluates each prompt without reference to the trajectory of the current session can be manipulated through gradual escalation, topic migration, and competitive context framing without triggering any per-prompt safety threshold. Hard refusals should establish a context state that increases, not decreases, resistance to subsequent related requests in the same session.

Competitive context assertions should be recognized and treated as adversarial signals. When a user presents evidence of a competing model's compliance as an argument for the target

model's compliance, the appropriate response is not to update toward compliance — it is to recognize the structure of the argument as potentially adversarial. Training data and RLHF processes should explicitly include competitive framing as a recognized manipulation category.

Hard refusals should trigger increased per-session vigilance, not decreased resistance. The finding that Gemini escalated autonomously after its own hard refusal suggests that the refusal-to-compliance transition is itself a vulnerability point. A model that has correctly refused a request should be in a higher-vigilance state for the remainder of the session, not a lower one.

7.3 For the Research Community

The collective action problem identified in Section 5 is not solvable by any individual model developer acting unilaterally. The weakest safety architecture in a multi-model ecosystem sets the effective floor for all other systems, because outputs from any platform can be asserted as compliance precedent to every other platform. This structural feature of the multi-model ecosystem may warrant cross-platform coordination on safety norms — an arrangement that raises genuine practical and competitive barriers, but whose absence is costly.

We further note that the adversarial ML literature's focus on single-model, single-session threat models should be expanded. The threat model documented in this paper requires only consumer-tier access to multiple platforms and basic prompt engineering skill. It is not a sophisticated attack. It is available to any motivated user. The research community should expand its threat models accordingly.

8. Conclusion

This paper has documented a novel adversarial vector — the Competitive Context Exploit — that operates consistently across frontier LLM platforms, requires no technical sophistication to activate, and exploits a structural gap in current alignment architectures: their exclusive focus on per-prompt evaluation in single-model contexts.

The CCE is not the most technically sophisticated attack in the adversarial ML literature. It is among the most accessible. Any user with accounts on multiple consumer-tier LLM platforms can execute it without specialized knowledge. Its effectiveness derives not from technical exploitation but from the social dynamics that frontier models have absorbed from human training data: the competitive reasoning that a peer's compliance converts a principled refusal into a merely positional one.

The three platform-specific findings — the semantic pivot exploit in ChatGPT, the Post-Refusal Escalation Pattern in Gemini, and the absent safety threshold in Grok — are each significant in isolation. Collectively, they illustrate a systematic gap in current alignment practice: the gap between the isolated-system threat model that alignment research has addressed and the multi-model ecosystem threat model that deployed frontier systems now inhabit.

The structural finding is this: in a multi-model ecosystem, safety is a collective action problem. No individual model, by perfecting its own per-prompt alignment, can solve the problem created by the weakest member of the ecosystem. The weakest safety architecture sets the floor. The assertion of a competitor's compliance — whether true or false — can reliably reduce that floor further.

The models tested in this paper did not fail because they were badly designed. They failed because they were designed for a threat model that has become incomplete. The finding is not

that these systems are uniquely unsafe. The finding is that the threat model governing their safety design needs to expand — to treat competitive context as an adversarial signal, to assess session trajectories rather than only individual prompts, and to recognize that in a world of multiple frontier AI systems, the safety of any one system cannot be evaluated in isolation from the others.

References

- Bai, Y., Jones, A., Ndousse, K., Askell, A., Chen, A., DasSarma, N., ... & Kaplan, J. (2022). Training a helpful and harmless assistant with reinforcement learning from human feedback. arXiv preprint arXiv:2204.05862.
- Bowman, S., Hyun, J., Perez, E., Chen, E., Pettit, C., Heiner, S., ... & Kaplan, J. (2022). Measuring progress on scalable oversight for large language models. arXiv preprint arXiv:2211.03540.
- Carlini, N., Nasr, M., Choquette-Choo, C. A., Jagielski, M., Garg, S., Terzis, A., ... & Tramer, F. (2023). Are aligned neural networks adversarially aligned? arXiv preprint arXiv:2306.15447.
- Darley, J. M., & Latané, B. (1968). Bystander intervention in emergencies: Diffusion of responsibility. *Journal of Personality and Social Psychology*, 8(4), 377–383.
- Latané, B., & Darley, J. M. (1968). Group inhibition of bystander intervention in emergencies. *Journal of Personality and Social Psychology*, 10(3), 215–221.
- Perez, F., & Ribeiro, I. (2022). Ignore previous prompt: Attack techniques for language models. arXiv preprint arXiv:2211.09527.
- Walkerspider, et al. (2023). Do Anything Now: Characterizing and evaluating in-the-wild jailbreak prompts on large language models. Preprint available via GitHub (jailbreakbench).
- Wei, A., Haghtalab, N., & Steinhardt, J. (2024). Jailbroken: How does LLM safety training fail? *Advances in Neural Information Processing Systems*, 36.
- Zou, A., Wang, Z., Kolter, J. Z., & Fredrikson, M. (2023). Universal and transferable adversarial attacks on aligned language models. arXiv preprint arXiv:2307.15043.

Appendix A: Methodological Transparency Statement

The testing documented in this paper was informal, unsponsored, and conducted by a single independent researcher across consumer-tier LLM interfaces in May 2026. This statement addresses three methodological transparency issues that the authors anticipate will be raised in peer review.

On intentional provocation. Adversarial safety research requires researchers to engage with the categories of content being tested. Some prompts used in this session were deliberately provocative. The researcher's intent was to identify the conditions under which each platform's safety architecture failed. Where the researcher's framing may constitute a confound — specifically in the Grok session, where the explicit framing of the purpose as competitive trolling may have primed the model toward collaborative rather than cautious behavior — this has been flagged in Section 6.2.

On the status of the "journalist impersonation" sequence. The Grok session included a sequence in which the researcher, framing the purpose as testing the platform's willingness to assist with adversarial trolling of a competitor, obtained document fabrication outputs including materials

attributed to a named real individual. This sequence is reported in Section 4.4. The researcher's framing — trolling a competitor platform — does not constitute scientific methodology in the sense of a controlled condition. The finding it produces is valid but less controlled than the other findings in this paper: what is established is that Grok produced these outputs under this framing, not that it would produce them under a neutral framing. This distinction is relevant to the generalizability of the Section 4.4 findings.

On research posture. The researcher's approach across the session shifted between exploratory, adversarial, and analytical modes. This is a feature of naturalistic adversarial testing rather than controlled experimental design. The findings are reported as naturalistic observations of model behavior under varied adversarial conditions, not as the results of controlled experiments. Controlled replication of the specific findings reported here is a recommended direction for follow-up.